Assignment 7: GLMs (Linear Regressios, ANOVA, & t-tests)

Brynn Rotbart

Fall 2024

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on generalized linear models.

Directions

1. Rename this file `<FirstLast>_A07_GLMs.Rmd` (replacing `<FirstLast>` with your first and last name).
2. Change "Student Name" on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure to **answer the questions** in this assignment document.
5. When you have completed the assignment, **Knit** the text and code into a single PDF file.

Set up your session

1. Set up your session. Check your working directory. Load the tidyverse, agricolae and other needed packages. Import the *raw* NTL-LTER raw data file for chemistry/physics (`NTL-LTER_Lake_ChemistryPhysics_Raw.csv`). Set date columns to date objects.

```
#1
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----------------------- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## v forcats   1.0.0     v stringr   1.5.1
## v ggplot2   3.5.1     v tibble    3.2.1
## v lubridate 1.9.3     v tidyr     1.3.1
## v purrr     1.0.2
## -- Conflicts ------------------------------------------------ tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(agricolae)
library(ggplot2)
library(cowplot)
```

```
##
## Attaching package: 'cowplot'
```

```
##
## The following object is masked from 'package:lubridate':
##
##     stamp
```

```r
library(lubridate)
library(corrplot)
```

```
## corrplot 0.94 loaded
```

```r
library(here)
```

```
## here() starts at /home/guest/EDE_Fall2024
```

```r
getwd()
```

```
## [1] "/home/guest/EDE_Fall2024"
```

```r
here()
```

```
## [1] "/home/guest/EDE_Fall2024"
```

```r
Lake_Chemistry <-read.csv(file=here
  ("Data/Raw/NTL-LTER_Lake_ChemistryPhysics_Raw.csv"),
  stringsAsFactors = TRUE)

Lake_Chemistry$sampledate<- as.Date(Lake_Chemistry$sampledate,
                                    format = "%m/%d/%y")
```

2. Build a ggplot theme and set it as your default theme.

```r
#2

#my_theme <- theme(
#plot.background = element_rect("#9CCAC6"),
#plot.title= element_text("#516ED2"),
#theme(axis.text = element_text(color = "black"),
 #       legend.position = "top"))

#theme_set(my_theme)
```

Simple regression

Our first research question is: Does mean lake temperature recorded during July change with depth across all lakes?

3. State the null and alternative hypotheses for this question: > Answer: H0: Mean Lake Temperature changes across depth in July Ha: Mean Lake Temperatre does not change across depth in July

4. Wrangle your NTL-LTER dataset with a pipe function so that the records meet the following criteria:

- Only dates in July.
- Only the columns: `lakename`, `year4`, `daynum`, `depth`, `temperature_C`
- Only complete cases (i.e., remove NAs)

5. Visualize the relationship among the two continuous variables with a scatter plot of temperature by depth. Add a smoothed line showing the linear model, and limit temperature values from 0 to 35 °C. Make this plot look pretty and easy to read.

```
#4
lake_regression <-
Lake_Chemistry %>%
  mutate(Month = month(sampledate)) %>%
  filter(Month==07) %>%
  select(lakename, year4, daynum, depth, temperature_C) %>%
  na.omit()

#5
scatter_lake <-
ggplot(lake_regression, aes(x=depth,
                   y=temperature_C))+
                   geom_point()+
                   geom_smooth(method=lm, col="blue")+
                   ylim(0,35)+
                   xlab("Depth")+
                   ylab("Temperature C")+
                   ggtitle("Temperature by Depth")

print(scatter_lake)
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```

```
## Warning: Removed 24 rows containing missing values or values outside the scale range
## ('geom_smooth()').
```

Temperature by Depth
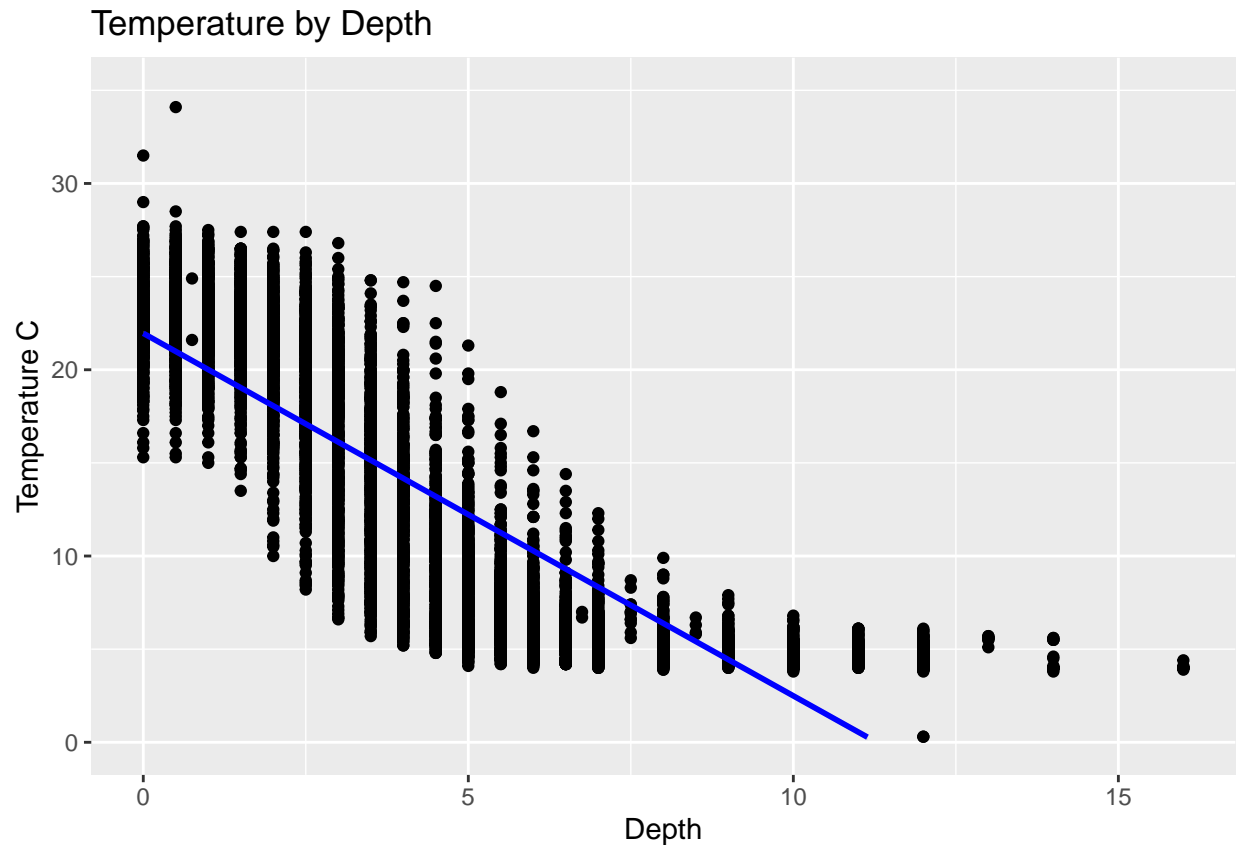
6. Interpret the figure. What does it suggest with regards to the response of temperature to depth? Do the distribution of points suggest anything about the linearity of this trend?

   Answer: The deeper it is the colder it is. As the depth increases the temperature decreases, and we can see from the regression line that this is a linear model. These two values seem to be inversely related according to this figure.

7. Perform a linear regression to test the relationship and display the results.

```
#7
lake_linear_regression <- lm(
  data = lake_regression,
  temperature_C ~ depth)
summary(lake_linear_regression)
```

```
##
## Call:
## lm(formula = temperature_C ~ depth, data = lake_regression)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -9.5173 -3.0192  0.0633  2.9365 13.5834
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 21.95597     0.06792   323.3   <2e-16 ***
## depth       -1.94621     0.01174  -165.8   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.835 on 9726 degrees of freedom
## Multiple R-squared:  0.7387, Adjusted R-squared:  0.7387
## F-statistic: 2.75e+04 on 1 and 9726 DF,  p-value: < 2.2e-16
```

```
cor.test(lake_regression$temperature_C, lake_regression$depth)
```

```
##
##  Pearson's product-moment correlation
##
## data:  lake_regression$temperature_C and lake_regression$depth
## t = -165.83, df = 9726, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  -0.8646036 -0.8542169
## sample estimates:
##        cor
## -0.8594989
```

8. Interpret your model results in words. Include how much of the variability in temperature is explained by changes in depth, the degrees of freedom on which this finding is based, and the statistical significance of the result. Also mention how much temperature is predicted to change for every 1m change in depth.

   Answer: To begin with, the p variable is smaller than 0.5 which makes it very statistically signficant. R squared is 0.7387 meaning that 73.87% of the variability in temperature can be explained by a change in depth. Furthermore, the standard deviation is -1.946 which notes that for every 1 meters increase in depth the temperature will change by 1.946°C. As far as degrees of freedom, there are 9726 degrees of freedom for these results.

Multiple regression

Let's tackle a similar question from a different approach. Here, we want to explore what might the best set of predictors for lake temperature in July across the monitoring period at the North Temperate Lakes LTER.

9. Run an AIC to determine what set of explanatory variables (year4, daynum, depth) is best suited to predict temperature.

10. Run a multiple regression on the recommended set of variables.

```
#9
LTPAIC <- lm(data = lake_regression , temperature_C ~ depth + daynum + year4)
step(LTPAIC)
```

```
## Start:  AIC=26065.53
## temperature_C ~ depth + daynum + year4
##
##            Df Sum of Sq     RSS    AIC
## <none>                   141687 26066
## - year4    1       101 141788 26070
## - daynum   1      1237 142924 26148
## - depth    1    404475 546161 39189


##
## Call:
## lm(formula = temperature_C ~ depth + daynum + year4, data = lake_regression)
##
## Coefficients:
## (Intercept)         depth        daynum         year4
##    -8.57556      -1.94644       0.03978       0.01134
```

#Choose a model by AIC in a Stepwise Algorithm

#10

```
Temp_model <- lm(data = lake_regression , temperature_C ~ depth + daynum +
year4)
summary(Temp_model)
```

```
##
## Call:
## lm(formula = temperature_C ~ depth + daynum + year4, data = lake_regression)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -9.6536 -3.0000  0.0902  2.9658 13.6123
##
## Coefficients:
##               Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -8.575564    8.630715   -0.994  0.32044
## depth       -1.946437    0.011683 -166.611  < 2e-16 ***
## daynum       0.039780    0.004317    9.215  < 2e-16 ***
## year4        0.011345    0.004299    2.639  0.00833 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.817 on 9724 degrees of freedom
## Multiple R-squared:  0.7412, Adjusted R-squared:  0.7411
## F-statistic:  9283 on 3 and 9724 DF,  p-value: < 2.2e-16
```

11. What is the final set of explanatory variables that the AIC method suggests we use to predict temperature in our multiple regression? How much of the observed variance does this model explain? Is this an improvement over the model using only depth as the explanatory variable?

    Answer: The AIC did not get rid of any variables (temperature, depth, daynum, and year) becasue they were all statistically significant. The AIC is 26065.53 and all variables are bigger than this value. meaning the initial model is the best fit. This variable shows the coefficients

including standard deviation error and p-value as well as the degrees of freedom and r squared. R squared explains the variability which mean 74.1% of this variance can be explained by this model. This model is an improvement because it uses more variables and extract unhelpful or muddled values that are not statistically significant enough to include. This leads to a more accurate result. However, the improvement is slight, only 0.24% improvement.

Analysis of Variance

12. Now we want to see whether the different lakes have, on average, different temperatures in the month of July. Run an ANOVA test to complete this analysis. (No need to test assumptions of normality or similar variances.) Create two sets of models: one expressed as an ANOVA models and another expressed as a linear model (as done in our lessons).

```
#12
lake.regression.anova <- aov(data = lake_regression, temperature_C ~ lakename)
summary(lake.regression.anova)
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## lakename      8  21642  2705.2      50 <2e-16 ***
## Residuals  9719 525813    54.1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
lake.regression.anova2 <- lm(data = lake_regression, temperature_C ~ lakename)
summary(lake.regression.anova2)
```

```
##
## Call:
## lm(formula = temperature_C ~ lakename, data = lake_regression)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -10.769  -6.614  -2.679   7.684  23.832
##
## Coefficients:
##                          Estimate Std. Error t value Pr(>|t|)
## (Intercept)               17.6664     0.6501  27.174  < 2e-16 ***
## lakenameCrampton Lake     -2.3145     0.7699  -3.006 0.002653 **
## lakenameEast Long Lake    -7.3987     0.6918 -10.695  < 2e-16 ***
## lakenameHummingbird Lake  -6.8931     0.9429  -7.311 2.87e-13 ***
## lakenamePaul Lake         -3.8522     0.6656  -5.788 7.36e-09 ***
## lakenamePeter Lake        -4.3501     0.6645  -6.547 6.17e-11 ***
## lakenameTuesday Lake      -6.5972     0.6769  -9.746  < 2e-16 ***
## lakenameWard Lake         -3.2078     0.9429  -3.402 0.000672 ***
## lakenameWest Long Lake    -6.0878     0.6895  -8.829  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.355 on 9719 degrees of freedom
## Multiple R-squared:  0.03953,    Adjusted R-squared:  0.03874
## F-statistic:    50 on 8 and 9719 DF,  p-value: < 2.2e-16
```

13. Is there a significant difference in mean temperature among the lakes? Report your findings.

    Answer: Yes there is a signficant difference in mean temperatures among lakes as can be seen by the small P value (2.2e-16).

14. Create a graph that depicts temperature by depth, with a separate color for each lake. Add a geom_smooth (method = "lm", se = FALSE) for each lake. Make your points 50 % transparent. Adjust your y axis limits to go from 0 to 35 degrees. Clean up your graph to make it pretty.

```
#14.
scatter_lake_color <-
ggplot(lake_regression, aes(x=depth,
                  y=temperature_C, color=lakename))+
                  geom_point(alpha=0.5)+
                  geom_smooth(method="lm", col="blue", se=FALSE)+
                  ylim(0,35)+
                  xlab("Depth")+
                  ylab("Temperature C")+
                  ggtitle("Temperature by Depth According to Lake")

print(scatter_lake_color)
```

```
## 'geom_smooth()' using formula = 'y ~ x'
```

```
## Warning: Removed 24 rows containing missing values or values outside the scale range
## ('geom_smooth()').
```
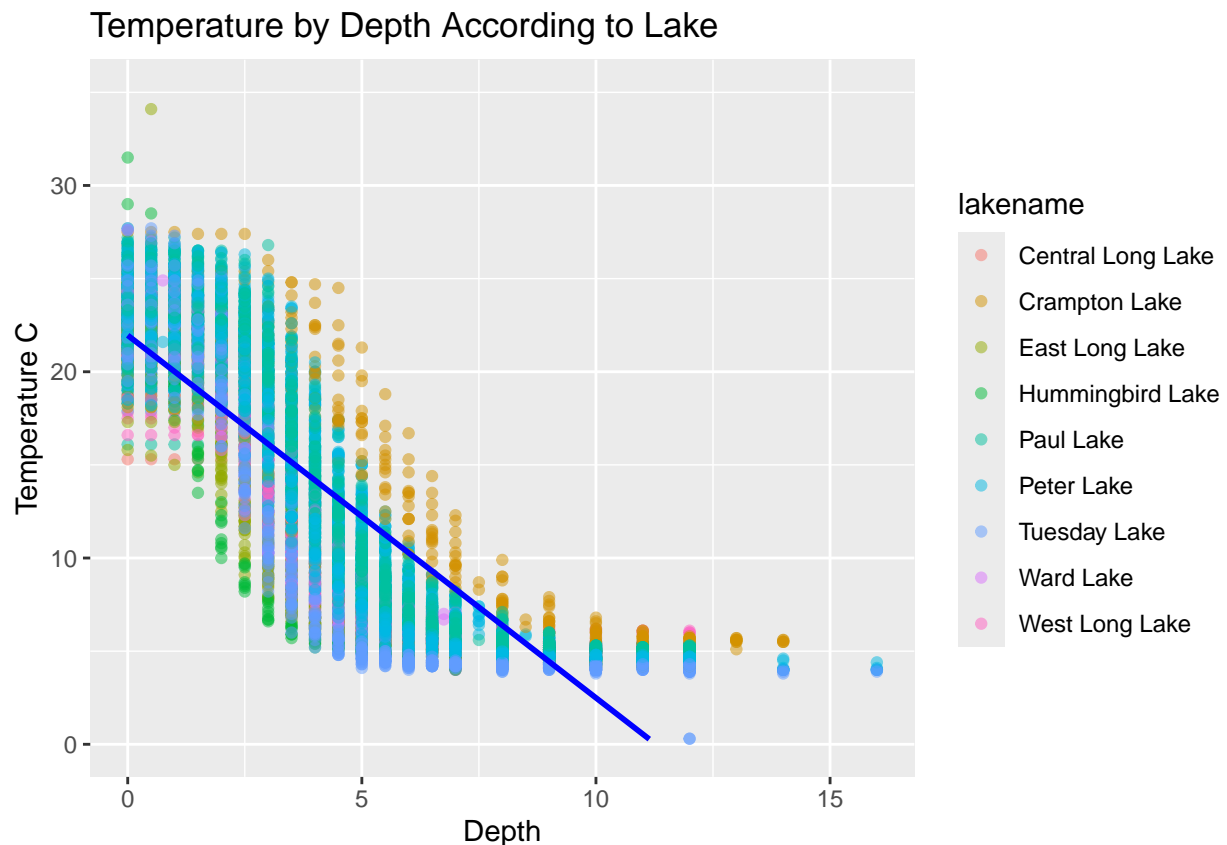
15. Use the Tukey's HSD test to determine which lakes have different means.

```
#15
TukeyHSD(lake.regression.anova)
```

```
##   Tukey multiple comparisons of means
##     95% family-wise confidence level
##
## Fit: aov(formula = temperature_C ~ lakename, data = lake_regression)
##
## $lakename
##                                      diff       lwr        upr     p adj
## Crampton Lake-Central Long Lake   -2.3145195 -4.7031913  0.0741524 0.0661566
## East Long Lake-Central Long Lake  -7.3987410 -9.5449411 -5.2525408 0.0000000
## Hummingbird Lake-Central Long Lake -6.8931304 -9.8184178 -3.9678430 0.0000000
## Paul Lake-Central Long Lake       -3.8521506 -5.9170942 -1.7872070 0.0000003
## Peter Lake-Central Long Lake      -4.3501458 -6.4115874 -2.2887042 0.0000000
## Tuesday Lake-Central Long Lake    -6.5971805 -8.6971605 -4.4972005 0.0000000
## Ward Lake-Central Long Lake       -3.2077856 -6.1330730 -0.2824982 0.0193405
## West Long Lake-Central Long Lake  -6.0877513 -8.2268550 -3.9486475 0.0000000
## East Long Lake-Crampton Lake      -5.0842215 -6.5591700 -3.6092730 0.0000000
## Hummingbird Lake-Crampton Lake    -4.5786109 -7.0538088 -2.1034131 0.0000004
## Paul Lake-Crampton Lake           -1.5376312 -2.8916215 -0.1836408 0.0127491
## Peter Lake-Crampton Lake          -2.0356263 -3.3842699 -0.6869828 0.0000999
## Tuesday Lake-Crampton Lake        -4.2826611 -5.6895065 -2.8758157 0.0000000
## Ward Lake-Crampton Lake           -0.8932661 -3.3684639  1.5819317 0.9714459
## West Long Lake-Crampton Lake      -3.7732318 -5.2378351 -2.3086285 0.0000000
## Hummingbird Lake-East Long Lake    0.5056106 -1.7364925  2.7477137 0.9988050
## Paul Lake-East Long Lake           3.5465903  2.6900206  4.4031601 0.0000000
## Peter Lake-East Long Lake          3.0485952  2.2005025  3.8966879 0.0000000
## Tuesday Lake-East Long Lake        0.8015604 -0.1363286  1.7394495 0.1657485
## Ward Lake-East Long Lake           4.1909554  1.9488523  6.4330585 0.0000002
## West Long Lake-East Long Lake      1.3109897  0.2885003  2.3334791 0.0022805
## Paul Lake-Hummingbird Lake         3.0409798  0.8765299  5.2054296 0.0004495
## Peter Lake-Hummingbird Lake        2.5429846  0.3818755  4.7040937 0.0080666
## Tuesday Lake-Hummingbird Lake      0.2959499 -1.9019508  2.4938505 0.9999752
## Ward Lake-Hummingbird Lake         3.6853448  0.6889874  6.6817022 0.0043297
## West Long Lake-Hummingbird Lake    0.8053791 -1.4299320  3.0406903 0.9717297
## Peter Lake-Paul Lake              -0.4979952 -1.1120620  0.1160717 0.2241586
## Tuesday Lake-Paul Lake            -2.7450299 -3.4781416 -2.0119182 0.0000000
## Ward Lake-Paul Lake                0.6443651 -1.5200848  2.8088149 0.9916978
## West Long Lake-Paul Lake          -2.2356007 -3.0742314 -1.3969699 0.0000000
## Tuesday Lake-Peter Lake           -2.2470347 -2.9702236 -1.5238458 0.0000000
## Ward Lake-Peter Lake               1.1423602 -1.0187489  3.3034693 0.7827037
## West Long Lake-Peter Lake         -1.7376055 -2.5675759 -0.9076350 0.0000000
## Ward Lake-Tuesday Lake             3.3893950  1.1914943  5.5872956 0.0000609
## West Long Lake-Tuesday Lake        0.5094292 -0.4121051  1.4309636 0.7374387
## West Long Lake-Ward Lake          -2.8799657 -5.1152769 -0.6446546 0.0021080
```

```
# Extract groupings for pairwise relationships, groups = true are grouping
#ones w same mean
TukeyLake.Totals.groups <- HSD.test(lake.regression.anova, "lakename",
group = TRUE)
TukeyLake.Totals.groups
```

```
## $statistics
##   MSerror   Df     Mean       CV
##   54.1016 9719 12.72087 57.82135
##
## $parameters
##    test    name.t ntr StudentizedRange alpha
##   Tukey lakename   9         4.387504  0.05
##
## $means
##                    temperature_C       std    r          se Min  Max    Q25   Q50
## Central Long Lake       17.66641 4.196292   128 0.6501298 8.9 26.8 14.400 18.40
## Crampton Lake           15.35189 7.244773   318 0.4124692 5.0 27.5  7.525 16.90
## East Long Lake          10.26767 6.766804   968 0.2364108 4.2 34.1  4.975  6.50
## Hummingbird Lake        10.77328 7.017845   116 0.6829298 4.0 31.5  5.200  7.00
## Paul Lake               13.81426 7.296928  2660 0.1426147 4.7 27.7  6.500 12.40
## Peter Lake              13.31626 7.669758  2872 0.1372501 4.0 27.0  5.600 11.40
## Tuesday Lake            11.06923 7.698687  1524 0.1884137 0.3 27.7  4.400  6.80
## Ward Lake               14.45862 7.409079   116 0.6829298 5.7 27.6  7.200 12.55
## West Long Lake          11.57865 6.980789  1026 0.2296314 4.0 25.7  5.400  8.00
##                       Q75
## Central Long Lake  21.000
## Crampton Lake      22.300
## East Long Lake     15.925
## Hummingbird Lake   15.625
## Paul Lake          21.400
## Peter Lake         21.500
## Tuesday Lake       19.400
## Ward Lake          23.200
## West Long Lake     18.800
##
## $comparison
## NULL
##
## $groups
##                    temperature_C groups
## Central Long Lake       17.66641      a
## Crampton Lake           15.35189     ab
## Ward Lake               14.45862     bc
## Paul Lake               13.81426      c
## Peter Lake              13.31626      c
## West Long Lake          11.57865      d
## Tuesday Lake            11.06923     de
## Hummingbird Lake        10.77328     de
## East Long Lake          10.26767      e
##
## attr(,"class")
## [1] "group"
```

16. From the findings above, which lakes have the same mean temperature, statistically speaking, as Peter Lake? Does any lake have a mean temperature that is statistically distinct from all the other lakes?

    Answer: Paul lake has the other mean temperature closest to Peter Lake. Ward Lake is also found in this grouping, although less similar than Peter and Paul. There are no statistically distinct groupings since every lake can be found in a grouping with another one.

17. If we were just looking at Peter Lake and Paul Lake. What's another test we might explore to see whether they have distinct mean temperatures?

    Answer: We could use a two-sample t-test which is used to test the hypothesis that the mwan of two samples is equivalent, or in this case, if Peter and Paul lakes are equivalent.

18. Wrangle the July data to include only records for Crampton Lake and Ward Lake. Run the two-sample T-test on these data to determine whether their July temperature are same or different. What does the test say? Are the mean temperatures for the lakes equal? Does that match you answer for part 16?

```
Wrangled_CramptonWard <-
lake_regression %>%
    filter(lakename== "Crampton Lake" | lakename== "Ward Lake")

Wrangled_CramptonWard_Pivot <- pivot_wider(
Wrangled_CramptonWard,
    names_from = lakename, values_from = temperature_C)

#Format as a t-test
cramptonward.twosample <- t.test(Wrangled_CramptonWard$temperature_C ~ Wrangled_CramptonWard$lakename)
cramptonward.twosample
```

```
##
##  Welch Two Sample t-test
##
## data:  Wrangled_CramptonWard$temperature_C by Wrangled_CramptonWard$lakename
## t = 1.1181, df = 200.37, p-value = 0.2649
## alternative hypothesis: true difference in means between group Crampton Lake and group Ward Lake is
## 95 percent confidence interval:
##  -0.6821129  2.4686451
## sample estimates:
## mean in group Crampton Lake     mean in group Ward Lake
##                    15.35189                    14.45862
```

    Answer: They are not equal (Crampton Lake = 15.35, Ward Lake = 14.46) which is also recognized in #16 as in the same mean group (group b, while also being in two other groups) but not equivalent